

What three markets need, and what three platforms do

Loupe, Maltego and Kaseware against criminal defence, corporate intelligence, and litigation support and disputes. Capability findings are drawn from published product documentation, release notes and pricing, current to 3 August 2026.

The three markets

Criminal defence

You receive the evidence. The prosecution hands you discovery — phone extractions, documents, bank records, bodycam, jail calls — and you did not collect any of it and cannot. The job is to know that material better than the other side, and to find what breaks their theory. Volumes are punishing: ten phones, six-figure document counts, court deadlines. Every assertion you make gets challenged, so every assertion has to point at a page. And what wins is frequently what is *absent* — the message that puts someone somewhere else, the date that does not line up.

What this buyer needs. Everything held in one model, so a person is the same person across the phone, the documents and the bank statements. The phone extraction treated as real records — individual calls and messages you can query — rather than a spreadsheet dump. A chronology that survives a courtroom. Every fact carrying its source. And something that actively hunts contradictions.

What this buyer does not need. To collect anything from the open web. To produce documents to another party. Machinery for batching a review team. These are genuinely irrelevant here, not weaknesses to excuse.

Corporate intelligence and business investigations

A client or their counsel asks a question: is this counterparty what they claim, where did the money go, who really owns this, was there fraud. The engagements are asset tracing, fraud and corruption, due diligence and disputes intelligence. The process is consistent across the field — scope it, plan your sources, collect, verify, analyse, write the report — and the report has to survive court.

What this buyer needs, in two distinct halves. At the front: reach outside the engagement, into corporate registries, beneficial ownership, sanctions and PEP lists, court filings, property deeds and press archives — and the ability to start from a name and pivot outward. At the back: take all of that plus whatever the client handed over, corroborate every finding against independent sources, hunt the inconsistencies and red flags, and produce a report where each claim is sourced. Bank records are the starting point of nearly every asset trace.

Why the split matters commercially. The front is where the market's money already goes, to data vendors the firm licenses anyway. The back is where the billable analyst hours go — and those hours are currently spent in Excel and Word. Junior consultants at these firms bill \$300–\$500 an hour; principals bill \$800–\$1,500. A six-week engagement with a team of four generates \$150,000–\$250,000 in fees.

Litigation support and disputes

Forensic accounting, fraud claims, asset tracing for enforcement, contentious insolvency, damages. You serve law firms, and your output is an expert report that gets cross-examined.

What this buyer needs. Everything the back half of corporate intelligence needs, held to a higher standard because someone will attack your method in a hearing. On top of that, things defence work never asks for: producing material to another party — redaction, privilege logs, numbered exhibits, load files — batching a review team through a large disclosure set, and being able to show your parsing has been independently validated. And because the client is a law firm, the firm-maturity requirements are real: proven at volume, plugs into the systems the firm already runs, and will still exist in three years.

One distinction that decides everything here. The industry that calls itself litigation support is court reporting, interpreting, record retrieval, transcription and trial graphics — bought by legal assistants and paralegals, won on responsiveness, serving personal injury, insurance and medical malpractice matters. That is a logistics and services business, and a rear-end collision file contains no evidence-modelling problem. The tier described above — disputes investigation and forensic accounting — is a different business wearing the same name.

The three platforms

Loupe

What it does. It takes everything that arrives in a matter — phone extractions, documents, bank statements, audio, video, images — and resolves it into a single model where the same person is the same object across all of it.

Underneath that: a phone extraction becomes first-class structured records, so individual calls, messages, contacts and locations are each queryable. Every extracted fact carries its verbatim quote, page and file, stored with the fact. Questions whose answer depends on relationships get answered across everything ingested, with a definite result — not a pile of documents that looked relevant. Names and places spoken in audio or shown in video become the same objects as the ones in documents. Bank records get categorised, counterparties resolved to who they actually are, and flows traced between parties. Chronologies persist as work product with their citations intact, as do an investigator's own conclusions — curated collections, approved facts, built narratives. It hunts for what contradicts a theory. Its answers are cited and its AI output can be audited. It runs single-tenant or on-premises, so evidence never sits in a shared system.

What it cannot do. Reach outside the matter at all — no registries, no sanctions lists, no adverse media, no court records. Connect to anything else — no document management, no practice management, no legal research, no third-party data. Collect evidence at source. Produce to another party. Run a review team. Its file-format range is narrower than the processing specialists', its case-management workflow is thin, and it is a young company with no audited security credentials and no published volume record.

Maltego

What it does. Reach, and it is the best in the field at it. Around 120 integrated data providers, metered by credits. Start from a name, email, domain, phone number or crypto wallet and pivot outward to whatever relates to it. Mature visual link analysis for exploring those connections. Real-time monitoring of social and news streams. Genuinely good preservation of what it finds on the open web — automated screenshots, chain-of-custody logging, defensible capture of public profiles and pages. It integrates into things, including letting a customer expose their own proprietary data through an on-premises server. It is an established, security-certified company with hundreds of thousands of users, and it can be bought on a card from €3,000 a year.

What it cannot do. Ingest the matter. Import is CSV and Excel — you map columns to entities. A discovery PDF has to be processed into a table outside the product first, and for a phone extraction there is no path at all. Its graph is an analyst's working document rather than a persistent model of the case: it holds what the analyst pulled, and the underlying data stays back in the sources. It has no real financial-records capability. Its AI reads and reports on the graph, and whether it cites sources is not documented anywhere. There are no narrative or hypothesis constructs and no contradiction hunting; its case feature is encrypted storage with collaboration — a container, not a work product. Law firms and defence practices are not a named customer type anywhere in its materials.

Kaseware

What it does. It runs the investigation. Built by former FBI Special Agents who built Sentinel, the FBI's own case management system, and it shows: intake, incident management, records, document and evidence management, task management and assignment, dispatch, a public portal, administrative controls, reporting and dashboards. The case exists as a managed object with a lifecycle, which is the thing an investigating unit actually needs and the thing most analytical tools do not have.

Around that sits real analytical tooling rather than workflow alone — link analysis, timelines, entity recognition and geospatial analysis. Collaboration and permissions are a strength: compartmentalised access, need-to-know

segregation and user auditing. There is a mobile app, and APIs and integrations. It is SOC 2, CJIS and TX-RAMP certified, and sold to law enforcement, fusion centres, intelligence agencies, corporate security, finance, aviation, energy and pharmaceutical customers.

What it cannot do. Build the evidence model. Documents and evidence are *managed* — stored, tracked, permissioned — not modelled, and there is no path from a discovery PDF to structured facts. Phone extractions are not first-class records. It has no financial-records capability worth the name. Facts do not carry their quote and page, so there is no per-claim provenance to point at under challenge. Its AI is published as "AI Services" with no documented detail, and nothing establishes that its answers cite sources. There are no hypothesis constructs and no contradiction hunting, and its durable work-product objects are thin. Its public materials do not describe what data formats it ingests at all. Law firms and defence practices are not among its named customer types.

Loupe and Kaseware in detail

These two look more alike than they are, and the resemblance is worth taking seriously rather than waving away. Both hold a case as a real object. Both draw link charts, build timelines and plot locations on a map. Both do entity recognition. Both are sold as investigation platforms. Anyone comparing feature lists will conclude they are competitors.

The one difference that produces all the others

Kaseware organises the investigation. Loupe interprets the evidence.

Kaseware knows the file exists, who touched it, whose task it is and what stage the case is at. Loupe knows what is inside the file and what it means.

That is why the apparent overlap is thinner than it looks. The link chart, the timeline and the map are *presentation* — and both platforms can draw all three. The difference is what fills them. In Kaseware, an analyst fills them in by hand, from material they have already read and understood. In Loupe, ingestion fills them: the extraction, the documents and the bank records become the entities and the events, and the chart is a view of the model rather than a drawing an analyst maintains.

So the comparison is not chart against chart. It is *who does the reading*.

What Loupe does that Kaseware does not

- **Turns a phone extraction into records.** Individual calls, messages, contacts and locations, each queryable and attributed to a person. Kaseware has no documented ingestion path for a forensic extraction; its evidence module stores and tracks items.
- **Turns documents into facts.** Kaseware's document management stores, tags, permissions and searches a file. Loupe reads the contents and lifts out the entities and assertions inside it. The difference is a filing cabinet with good search, against a model of what the documents say.
- **Resolves one identity across different kinds of evidence.** The same person is the same object whether they appear in the phone, a contract or a bank statement. Kaseware's entity recognition operates inside its own records; nothing indicates it unifies an identity across an extraction, a statement and a PDF.
- **Treats financial records as analysable data** — categorising transactions, resolving who a counterparty actually is, and tracing flows between parties. For asset tracing this is not a feature, it is the job. Kaseware has essentially nothing here.
- **Attaches a source to every individual claim** — verbatim quote, page and file, stored with the fact. Kaseware maintains chain of custody on *items* of evidence. Under cross-examination those are different things: one proves the exhibit is what it says it is, the other proves the sentence in your report is supported.
- **Answers relationship questions across everything ingested**, with a definite result, rather than searching records and returning what matched.
- **Brings audio, video and images into the same model** — names and places spoken or shown become the same objects as those in documents.
- **Gives cited answers and auditable AI output.** Kaseware publishes "AI Services" with no documented detail and nothing establishing that its answers cite sources.
- **Hunts for what contradicts a theory.** Kaseware has nothing of the kind.

- **Holds investigator judgment as first-class objects** — curated collections, approved facts, built narratives — rather than case notes and generated reports.

What Kaseware does that Loupe does not

This list is longer, and most of it is real rather than cosmetic.

- **Runs a unit's caseload.** Case intake, incident management, assignment to named investigators, tasks with due dates, status and lifecycle, supervisor oversight, dispatch. Loupe's equivalent is thin.
- **Records management** as a system in its own right, not only case files.
- **A public portal** for intake from outside the organisation. Loupe has nothing comparable.
- **Compartmentalised access across a large organisation** — need-to-know segregation and user auditing at a level an agency or a corporate security function requires. Loupe's permissions are basic.
- **Management reporting and dashboards** on the caseload itself: volumes, status, workload. This is reporting about the work rather than about the evidence, and it is what a supervisor buys.
- **A mobile app** for field use.
- **APIs and integrations** into surrounding systems. Loupe has none at all.
- **SOC 2, CJIS and TX-RAMP.** Loupe holds none of these, and for an agency or a corporate security team they are frequently the first question rather than a later one.
- **Some open-source research tooling.** Modest, but not nothing, which is what Loupe has.
- **Physical and item-level evidence custody** — tracking evidence as property. Loupe models the contents of digital evidence; it is not a property-room system.
- **Institutional credibility.** Founded by the former FBI agents who built Sentinel, and sold across law enforcement, fusion centres, intelligence agencies, corporate security, finance, aviation, energy and pharmaceutical customers. Loupe is pre-revenue.

Where both are equally weak

Neither does document review and production in the eDiscovery sense: no batching work out to a review team, no coding panels, no reviewer quality control, no redaction, no privilege logs, no numbered exhibits or load files. In any matter where the deliverable is a production to another party, neither platform is the answer.

The commercial risk this creates

Both products get described as investigation platforms, and in corporate intelligence they can end up on the same shortlist and the same budget line. **In a comparison framed as investigation management, Kaseware wins nearly everything procurement scores** — credentials, integrations, permissions, reporting, mobile, vendor track record — **and loses nearly everything the analyst needs.** That framing is unfavourable and avoidable. The comparison to invite is what happens to the evidence after it arrives; the comparison to decline is which product manages a caseload better, because Kaseware manages a caseload better.

The genuinely useful reading is that a firm can run Kaseware and still have every problem Loupe solves, because organising a case and understanding its contents are different jobs. That is the sentence to lead with when the two come up together.

Who fits where

The structural picture

The three platforms answer three different questions about the same investigation:

- **Kaseware — is the work being run properly?** Who is assigned it, what stage is it at, who is allowed to see it, what does the supervisor need to know.
- **Maltego — what exists outside the case?** Registries, sanctions lists, adverse media, social material, and how it connects to a name you already have.
- **Loupe — what does the material actually say?** One model across every kind of evidence, each claim traceable to its source, a chronology that holds, and the facts that contradict the theory.

All three questions are legitimate and a serious engagement asks all three. What matters commercially is that **neither of the other two answers Loupe's question, and Loupe's two weakest areas are exactly what they are respectively good at** — running the work is Kaseware, reaching outward is Maltego. That is the honest shape of the field, and it is more useful than a ranking because it says what to build, what to integrate and what to leave alone.

Criminal defence

Loupe fits, and the other two are not in this market. What the buyer needs and what Loupe does are the same list — it satisfies every one of the twelve requirements this buyer has.

Maltego satisfies none of them. The evidence arrives from the prosecution and there is no way to get it into the product. Kaseware satisfies two, and is a case-management system for an investigating agency rather than an evidence platform for the receiving side — it is sold primarily to the other side of the table.

Corporate intelligence

This market splits cleanly down the middle of the engagement, and all three platforms sit in different parts of it. Maltego owns the sourcing and collecting. Kaseware owns running the engagement — genuinely useful here, because this market does need case workflow, need-to-know collaboration and enterprise-grade credentials, and it is stronger on all three than Loupe is. Loupe owns the corroborating, analysing and reporting.

This is also the market where Kaseware is most likely to be mistaken for a competitor, and where the framing of the comparison decides the outcome. See the detailed section above.

The halves are worth different amounts. The front is spend that already goes to licensed data vendors and is not really available to compete for — winning it means becoming a data business. The back is billable hours sitting in spreadsheets. Loupe's missing piece is reach, which can be integrated. Maltego's missing piece is the evidence model, which is architecture. Kaseware's missing piece is also the evidence model.

Of the nineteen things this buyer needs, Loupe meets twelve, Kaseware nine and Maltego seven — and the ones Loupe misses are reach, integrations, workflow and company maturity, not analysis.

Litigation support and disputes

The industry as it defines itself is not addressable, and should not be claimed. Neither is eDiscovery and managed review: it is the fastest-growing thing in the vicinity and precisely the capability set Loupe is weakest at —

no batching, no coding panels, no reviewer QC, no redaction, no privilege logs, no production formats. The incumbents own it outright.

The disputes-investigation tier is real, and it is the weakest of the three markets for all three platforms. Loupe brings the analysis and the provenance the expert report is built on, and is missing production mechanics, review machinery and track record — maturity items rather than architectural ones, and most already on the roadmap. Kaseware and Maltego are both further off, for the same reason in each case: neither can ingest the matter, and neither handles the financial records that disputes work actually runs on.

The right framing is that these firms serve law firms, so selling to them is selling *through* them. That makes it a channel worth raising in a meeting, and not worth a slide.

What is safe to say

The claim that carries all three markets

One finding does more work than everything else, and it is the one that survives being checked: **of the platforms examined, not one holds device data, financial records and documents in a single model**. That is not an artefact of how the comparison was drawn — it is structural. The companies that can model evidence sell their analytical tiers to law enforcement and intelligence, or they model devices and nothing else. The companies that will sell to this buyer are excellent at documents and treat a phone extraction as an attachment. Modelling evidence and serving this buyer are different engineering problems sold to different customers, so nobody has had a commercial reason to do both.

Underneath it sits the thing no competitor can contest: this was built because real casework needed it and nothing trialled was fit for purpose. Say "a number of platforms" and name none of them — naming one invites a comparison on their terms, and the audience almost certainly does not know them anyway.

Criminal defence — claim it fully

The only market where the product is complete as it stands, with no roadmap gap between what the buyer needs and what exists.

Safe without hedging: every feature came from a working investigator on live federal casework; cases have been worked to completion in the platform; three attorneys on those cases have access. This buyer receives evidence, cannot collect it, and cannot buy the analytics the other side runs on.

Not safe: that anything is signed, that it is proven at volume, or that it holds audited security credentials. The honest framing of the security gap is that it blocks the first external customer rather than making the product a bad fit — and it is the specific thing funding pays for. Say that plainly rather than letting someone find it.

Corporate intelligence — claim half of it

Safe: the corroborate–analyse–report half, which is where the billable hours sit, at \$300–\$500 an hour, currently going into Excel and Word. It is the same product as the beachhead, so pursuing it costs no roadmap.

Not safe: any open-source or external-data capability, because there is none. That means due diligence and KYC are not addressable — that half runs on licensed data the firms already buy, and competing there means becoming a different company. Conceding it early makes the half that is claimed more credible.

Two precision points. **"Corporate intelligence" means the investigative services firm** — the boutique or major-practice team doing fraud, corruption and asset tracing — not the in-house corporate investigations function, which is a workflow-and-procurement buyer and a materially worse fit. And **do not cite the EU due diligence directive as a tailwind**: vendor market reports do, and it has been pushed to 2028 transposition and 2029 application with narrowed scope. It fails the first check anyone runs.

Litigation support — sell through it, not to it

Claim the disputes-investigation tier and the channel. Do not claim the litigation support industry, and do not claim eDiscovery or managed review.

The demand driver worth citing is litigation funding: \$29.2 billion in 2026, with 42% of commercial disputes now carrying third-party funding. A funder is a cost-disciplined buyer with a direct interest in reducing the hours a matter consumes, which makes it the most receptive audience a labour-substitution argument can have.

On competitors, in any market

Answer with the category and that vendor's own published limit, not the name. An audience that has not heard of a vendor learns nothing from the name and everything from the limit. Where a competitor is genuinely strong, say so — the complement is more credible than the contest, and in two of these three cases it is also true.

Two regulatory facts worth more than any market-size number

Both are published by institutions rather than vendors, both are dateable, and neither requires the audience to have heard of a single competitor.

Proposed Federal Rule of Evidence 707. The Judicial Conference has recommended subjecting machine-generated evidence to the same reliability standard as expert testimony — the proponent must show the output rests on sufficient facts and data, produced by reliable methods, reliably applied. Public comment closed 16 February 2026; if approved it takes effect 1 December 2027.

Two limits, and they must be stated before this is used anywhere. The rule is not law: it has cleared comment, not the Standing Committee, Judicial Conference, Supreme Court and Congress. And it bites only where machine output is offered *without* an expert witness — in defence and disputes work there is usually a human testifying to the analysis, in which case Rule 702 already governs. So this is a direction-of-travel argument, not a compliance mandate, and claiming otherwise fails on contact with any litigator. Used correctly it is still strong: the courts are moving to make AI-derived findings challengeable on method, which rewards an architecture where every claim carries its verbatim quote, page and source, and makes a platform that cannot show its method a liability.

Practice Direction 57AD. Permanent in the Business and Property Courts of England and Wales since 1 October 2022. Extended Disclosure runs on models — known adverse documents only, key documents relied on plus known adverse, or particular documents and narrow classes — with the court holding parties to burden and cost being reasonable and proportionate. A regime built around targeted, proportionate, issue-led disclosure favours a platform that can find the specific document and prove where it came from, over one that reviews everything.