



Find the signal in everything.

**Everyone can cite a document. Nobody can query a case.**

An investigation platform for fraud and criminal casework — in production use on real federal matters.

Dublin, Ireland.

---

## THE PROBLEM

# Modern casework is a data problem wearing a legal costume.

---

A single serious-fraud or criminal case now arrives as:

- thousands of documents
- multi-gigabyte phone extractions, hundreds of thousands of events
- bank records spanning tens of thousands of transactions
- hours of recorded audio

The tools investigators have are **viewers**. One phone at a time. One PDF at a time. One spreadsheet at a time.

**And the tools that could hold it throw it away.** Relativity's own processing documentation says call logs emerge “only as call log data in Excel format, not as individual records,” and that multi-device extractions are unsupported. Everlaw splits conversations every thousand messages; Reveal slices them into 24-hour blocks.

The connections — the same person in a witness statement, a WhatsApp thread and a beneficiary field; the meeting two streets from the cash withdrawal — live in investigators' heads and on whiteboards.

**They don't scale, they don't survive staff turnover, and they can't be handed to a court.**

**We didn't start with a product idea. We started with a case we couldn't work.**

*We tried to buy this. Building it was the second choice, not the first.*

# Three curves crossed.

---

## 01

**Digital evidence volume exploded.** A phone extraction now appears in nearly every case.

## 02

**LLMs became good enough** to read evidence at scale.

## 03

**The category converged on cited answers — and stopped there.**

Relativity folded aiR into its base rate in November 2025; Everlaw bundled three AI features into core pricing the same month; DISCO announced agentic AI at no extra cost in February 2026. Everyone can cite a document. Nobody modelled the case.

That is the gap Loupe was built into:

**the layer underneath the answer — a model of the whole case, so a question can be put to all of it at once.**

Only ~25% of anti-fraud programmes use AI/ML today — and only **~6% of fraud examiners are confident explaining how their AI reaches its conclusions.**

## PRODUCT

# Evidence in. One connected model out.

Phone extractions and call logs, documents, images, audio, financial records — everything a case arrives with — resolved at ingestion into **one model**, where a call, a transfer and a line in a subpoena return are the same kind of citable object.

## ONE MODEL, MANY PERSPECTIVES

Financial explorer

Multi-phone view

Timeline

Graph

Agent query

### SCREENSHOT — A LOUPE

Passages, entities and events from several sources,  
bound into one evidenced finding.

### SCREENSHOT — CITED AI ANSWER, MID-CLICK

Opening the source document  
at the cited page.

**A Loupe is a bonded collection** — the passages, entities, events and transactions that establish one thing, each carrying the quote, page and file it came from.

Investigators work the perspectives, find what matters, and bind it into a Loupe. **Loupes move into the case, and the case is built out of them** — conclusion, justification and evidence, assembled as the investigation runs rather than written up at the end.

DIFFERENTIATION

# Why this isn't ChatGPT with files.

|                      | A straight LLM                                                                               | Loupe                                                                                                                                          |
|----------------------|----------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| Where the case lives | A context window, forgotten after                                                            | A persistent case knowledge graph — outlives every conversation                                                                                |
| Scale                | Hundreds of pages at best                                                                    | A real case: ten phones, 200,000 documents, 500 hours of audio — 435GB, all of it modelled                                                     |
| What the AI does     | Retrieves a sample of what fits and reasons over it in prose. You cannot know what it missed | Traverses a model of the whole case and returns the complete matching set — every claim cited, and the citation covers all of it, not a sample |
| Confidentiality      | Multi-tenant cloud                                                                           | Single-tenant: evidence never leaves the customer's instance                                                                                   |
| Human judgment       | Lives outside the tool — lost between sessions                                               | Highlights, significance and Loupes — bonded collections of evidence — are durable objects in the case model                                   |
| Model dependence     | The model is the product                                                                     | Models are interchangeable components. The moat is the evidence layer.                                                                         |

And not only ChatGPT — the specialists publish their own ceilings. Relativity caps at 300,000 documents per index; Reveal states its search is “designed for precision, not recall”.

A straight LLM gives you a well-written opinion about the fraction of the evidence that fits in its context window. Loupe models all of it — and every claim carries its receipt.

# Forty-seven platforms. The intersection is empty.

We assessed the category on thirty axes before we built. It splits into two halves that never meet — and the buyer sits in the gap between them.

## They model the evidence — and won't sell it to you

Cellebrite · Magnet · Palantir · Pathfinder

They own the extraction and build real graphs. Pathfinder is sold only to law enforcement and intelligence; GrayKey not to the private sector at all. What a practice can buy is examiner tooling — five kinds of device, no kinds of document.

## They'll sell to you — and treat a phone as an attachment

Relativity · Everlaw · DISCO · Reveal · Harvey · CoCounsel

Excellent at documents, purchasable, and the phone data does not survive the door. Call logs emerge as Excel, conversations are split every thousand messages, multi-device extractions are unsupported.

## Same buyer as us — and no case model

TrialKit · Matey · JusticeText · Longeye

Four funded companies selling AI to criminal defence, \$2.5–7.5M seed each, with the bar associations and the certifications. **None of them has a graph.** All four answer by retrieval — which returns what looked relevant and cannot tell you what it missed. **On a defence matter the thing you never saw is the thing that loses the case.**

**Not one of the forty-seven holds device data, financial records and documents in a single model** — and that is structural, not an accident of scoring. Modelling evidence and serving this buyer are different problems sold to different customers.

**We know because we went looking for a platform to run live federal casework, tried several, and none could work the case.**

## MARKET AND LEAD CUSTOMER

# Every feature was requested by a working investigator on a live case.

Our CPO is our design partner. Development has been customer-led from the first commit — not a roadmap we invented and then validated, but a backlog generated by real federal casework. Cases have been worked to completion in the platform, and **three attorneys on those cases have been given access to it.**

## Why this tier is empty — and stays empty.

Enterprise investigative-intelligence platforms are sold with configuration and support attached — the right model for agencies and large institutions, who arrive with budget, procurement and implementation resource.

The tier below has the same casework complexity and the same evidence volumes — and none of that budget. **Loupe is built for it directly: usable out of the box, by a practice with no implementation resource at all.**

**This buyer receives evidence, cannot collect it, and cannot buy the analytics the other side runs on.**

**That is the whole opportunity.**

## Beachhead

Boutique forensic and criminal-defence practices — 10–50 people, Big-Four-class casework without Big-Four tooling budgets.

## Expansion

Financial-crime & fraud units · law enforcement & prosecutors · law firms & disputes teams · corporate investigations & compliance.

## Channel

This buyer is organised and reachable. ACFE has ~95k members and ~60k CFEs worldwide, with chapters, conferences and directories — including an Ireland chapter since 2016. The defence bar is the same shape: a competitor reached national presence through a national association and one state bar.

## How many of them

**48,000–125,000 evidence-heavy private criminal-defence matters run in the US every year**, across a target population of one to five thousand practices. The UK adds ~80,200 open Crown Court cases, rising.

CUSTOMER VALUE

Today a serious case costs \$38–139k.  
With Loupe, \$16–20k.

Same case, same evidence, both columns. Ranges run from a two-phone fraud matter to a ten-phone federal case.

| Per serious case                                      | Today              | With Loupe        |
|-------------------------------------------------------|--------------------|-------------------|
| Forensic extraction — devices still go to a vendor    | \$3,650 – 8,450    | \$3,650 – 8,450   |
| Processing, hosting and examiner analysis             | \$34,200 – 130,575 | —                 |
| Loupe, per matter                                     | —                  | \$12,000          |
| Cash per case                                         | \$37,850 – 139,025 | \$15,650 – 20,450 |
| Investigator and paralegal hours on evidence handling | 58 – 100 hrs       | 29 – 50 hrs       |

**Saved: \$22,000–119,000 and 29–50 hours per case.** Across a practice running fifteen to thirty serious matters a year, that is **\$0.6–1.4M and 550–1,150 hours returned.**

The extraction bill is unchanged — we analyse what was produced, we do not acquire it. Everything downstream of it is what we replace.



## GROWTH PATH

# Ten million in ARR is 104 firms and under 2% of one market.

Priced at a blended \$12,000 per matter. A boutique running eight evidence-heavy matters a year is worth \$96,000 annually; an active practice at twenty is worth \$240,000. These figures use the conservative rate throughout.

### STAGE 1 • PROOF

Owl & Ireland

**\$0.5M ARR**

5 customers

42 matters / yr

References and a validated price. Not a revenue stage.

### STAGE 2 • THE ENGINE

United States

**\$5M ARR**

52 customers

417 matters / yr

0.3–0.9% of 48,000–125,000 evidence-heavy matters a year.

### STAGE 3 • SECOND GEOGRAPHY

+ United Kingdom

**\$10M ARR**

104 customers

833 matters / yr

80,200 open Crown Court cases, rising. Same buyer, same channel.

### STAGE 4 • EXPANSION

+ Europe

**\$25M ARR**

260 customers

2,083 matters / yr

\$2.41B forensics, \$6.81B legal tech. Partner-led per jurisdiction.

**Each stage is gated on one thing, which is what makes the sequence testable.** Stage 1: a price validated on real matters. Stage 2: the bar-association channel converting. Stage 3: the per-matter model rebuilt on UK vendor and Legal Aid Agency rates. Stage 4: partners, because every jurisdiction carries its own procedure and disclosure regime.

## TEAM

# We didn't research this market. We worked in it — from both sides.

---

## Neil Byrne

### CEO

Seven years at Siren, the Irish investigative-intelligence platform — the same category, selling to law enforcement, intelligence and financial-crime teams. **Solutions engineer: managed customer accounts and built the customisations each client needed to make the platform fit their process.** Spoke with hundreds of analysts across conference talks and deployments. Now a sales solutions engineer at UdeMy across large-enterprise and strategic accounts — the same discipline at enterprise scale. 10+ years in data engineering and AI systems.

**Full-time on close.**

## Conor Bowles

### CTO

**Architect of the case model.** The approach the platform rests on is his: extract entities and events from evidence with an LLM at ingestion, resolve them into a graph, and query that concrete structure afterwards — so an answer is a traversal over modelled evidence, not a retrieval over text. He wrote the first commit and has led engineering since. **Drives the enterprise platform**, while features are proven on live casework first and promoted once they mature — so nothing reaches the deployable product until it has survived a real matter.

**Full-time on close.**

## Alexandra Solórzano

### CPO

10+ years as a licensed private investigator on federal criminal defence and financial fraud casework — and the product's direction, daily. **Multiple calls a week; the backlog is generated from live matters as she works them**, which is why development has never run on a roadmap invented in advance. A grounded visionary: every feature traces to a real case, and she has run federal cases to completion inside the platform. US-based.

**The domain authority — and the customer the product was built for.**

Seven years building investigative intelligence software for enterprise buyers, and a decade running federal casework as the kind of investigator who has to live with the result. **Our CPO has run federal cases to completion inside this platform.**

## STATUS AND TRACTION

# Loupe is in production use on real federal casework.

---

Built inside **Owl Consultancy**, a working US private-investigations practice, and used there to carry federal cases **to completion** — multi-gigabyte phone extractions, tens of thousands of curated financial transactions, full document corpora. **Three attorneys on those cases have been given access to it.** Attorneys have retained the practice specifically because of the platform — unprompted, before any marketing existed.

**This is not a prototype seeking a first user. It is an application that works, in the hands of the investigators it was built for, on cases with real stakes.**

### First external pilot in motion.

A US private-investigations firm led by a former FBI undercover agent — clients across Miami, New York, California and the Midwest — has agreed to send its next large-document matter through Loupe as a paid trial at cost, and to put a mutual MSA in front of us. Sourced through casework, not marketing.

### Built so far, verifiable

- Federal cases carried to completion — multi-gigabyte phone extractions, tens of thousands of curated financial transactions, full document corpora
- A React console, a FastAPI case API and a separate asynchronous evidence engine, built since December 2025 by two engineers
- Neo4j (case graph) · Postgres (cases, auth, audit) · ChromaDB (retrieval) · Redis (jobs)
- Three AI providers behind one routing policy — the model is a swappable component
- Single-tenant deployment: one isolated stack per customer

### Release discipline

Onboarding an external customer's evidence runs behind a **12-stage security-gated release plan. No external case data enters the platform until the security, isolation, backup and legal stages pass.**

## THE ASK

# €100k. What it buys: the first paying customer outside the practice that built it.

---

The platform works and has closed real federal cases. What stands between it and external revenue is the **security gate** — the isolation, backup and legal stages that let another firm's privileged evidence enter the system — and the pilot itself.

### Use of funds

- **Two founders full-time** — clearing the security gate and running the agreed pilot to a signed contract. The pilot candidate is already secured.
- **The security evidence a firm's counsel asks for** — Cyber Essentials, an independent penetration test and a written security package, plus starting the SOC 2 observation window. Cheap and fast; full certification is bought when a deal requires it.
- **Single-tenant infrastructure** for the first external customers — one instance each, so it scales with customers rather than ahead of them.
- **Processing and inference** across pilot matters — the cost model the pilot establishes.

### What we want from NDRC beyond capital

- **Go-to-market guidance.** The channel is identified — bar associations and professional bodies — and pricing is per-matter rather than per-seat. Both need pressure-testing by people who have built a repeatable motion.
- Dublin base at Dogpatch Labs
- EIR time on pricing structure and enterprise sales motion
- Introductions into Irish and UK professional-services buyers — ACFE Ireland, Chartered Accountants Ireland forensic group, Law Society criminal law committee

**We are not asking anyone to fund the invention of a product. We are asking them to fund the distance between works and sold.**