



Find the signal in everything.

An investigation platform for fraud and criminal casework.

Dublin, Ireland.

THE PROBLEM

Modern casework is a data problem wearing a legal costume.

A single serious-fraud or criminal case now arrives as:

- thousands of documents
- multi-gigabyte phone extractions, hundreds of thousands of events
- bank records spanning tens of thousands of transactions
- hours of recorded audio

The tools investigators have are **viewers**. One phone at a time. One PDF at a time. One spreadsheet at a time.

The connections — the same person in a witness statement, a WhatsApp thread and a beneficiary field; the meeting two streets from the cash withdrawal — live in investigators' heads and on whiteboards.

They don't scale, they don't survive staff turnover, and they can't be handed to a court.

We didn't start with a product idea. We started with a case we couldn't work.

We tried to buy this. Building it was the second choice, not the first.

Three curves crossed.

01

Digital evidence volume exploded. A phone extraction now appears in nearly every case.

02

LLMs became good enough to read evidence at scale.

03

Bare LLMs remain unusable in evidence work — no provenance, no audit, no confidentiality.

The gap between (2) and (3) is exactly where Loupe sits:

the trust layer that makes AI admissible in investigative work.

Only ~25% of anti-fraud programmes use AI/ML today — and only **~6% of fraud examiners are confident explaining how their AI reaches its conclusions.**

ACFE 2026 anti-fraud technology benchmarking

PRODUCT

Evidence in. One connected model out.

SCREENSHOT — A LOUPE

Five highlighted passages from three documents,
bound into one evidenced narrative timeline.

SCREENSHOT — CITED AI ANSWER, MID-CLICK

Opening the source document
at the cited page.

Every fact carries the verbatim quote, source location and confidence it rests on.

Investigators highlight what matters and bind it into **Loupes** — evidenced collections that carry the narrative, not just the extraction. The product is named after them.

Ingest → Ground → Connect → Explore → Prove

DIFFERENTIATION

Why this isn't ChatGPT with files.

	A straight LLM	Loupe
Where the case lives	A context window, forgotten after	A persistent case knowledge graph — outlives every conversation
Scale	Hundreds of pages at best	Entire corpora: thousands of documents, multi-gigabyte phone reports
What the AI does	Answers in prose, one conversation at a time	Cites every claim to the page — and runs tools: builds Loupes, groupings and outputs from plain instruction
Confidentiality	Multi-tenant cloud	Single-tenant: evidence never leaves the customer's instance
Human judgment	Lives outside the tool — lost between sessions	Highlights, significance and Loupes — bonded collections of evidence — are durable objects in the case model
Model dependence	The model is the product	Models are interchangeable components. The moat is the evidence layer.

A straight LLM gives you a well-written opinion about the fraction of the evidence that fits in its context window. Loupe gives you a complete, permanent, cross-referenced model of all of it — where **every claim carries its receipt.**

MARKET AND LEAD CUSTOMER

Every feature was requested by a working investigator on a live case.

Our CPO is our design partner. Development has been customer-led from the first commit — not a roadmap we invented and then validated, but a backlog generated by real federal casework. Cases have been worked to completion in the platform, and **three attorneys on those cases have been given access to it.**

Why this tier is empty — and stays empty.

Enterprise investigative-intelligence platforms are sold with configuration and support attached — the right model for agencies and large institutions, who arrive with budget, procurement and implementation resource.

The tier below has the same casework complexity and the same evidence volumes — and none of that budget. **Loupe is built for it directly: usable out of the box, by a practice with no implementation resource at all.**

We did not set out to build a platform. We went looking for one, tried several, and none could work the case.

Beachhead

Boutique forensic and criminal-defence practices — 10–50 people, Big-Four-class casework without Big-Four tooling budgets.

Expansion

Financial-crime & fraud units · law enforcement & prosecutors · law firms & disputes teams · corporate investigations & compliance.

Channel

This buyer is organised and reachable. ACFE has ~95k members and ~60k CFEs worldwide, with chapters, conferences and directories — including an Ireland chapter since 2016. The defence bar is the same shape: a competitor reached national presence through a national association and one state bar.

Why we know the gap is real

We tried a number of platforms on live federal casework before we built anything. None of them could work the case.

GROWTH PATH

Ten million in ARR is 104 firms and under 2% of one market.

Priced at a blended \$12,000 per matter. A boutique running eight evidence-heavy matters a year is worth \$96,000 annually; an active practice at twenty is worth \$240,000. These figures use the conservative rate throughout.

STAGE 1 • PROOF

Owl & Ireland

\$0.5M ARR

5 customers

42 matters / yr

References and a validated price. Not a revenue stage.

STAGE 2 • THE ENGINE

United States

\$5M ARR

52 customers

417 matters / yr

0.3–0.9% of 48,000–125,000 evidence-heavy matters a year.

STAGE 3 • SECOND GEOGRAPHY

+ United Kingdom

\$10M ARR

104 customers

833 matters / yr

80,200 open Crown Court cases, rising. Same buyer, same channel.

STAGE 4 • EXPANSION

+ Europe

\$25M ARR

260 customers

2,083 matters / yr

\$2.41B forensics, \$6.81B legal tech. Partner-led per jurisdiction.

TEAM

We didn't research this market. We worked in it — from both sides.

Neil Byrne

CEO

Seven years at Siren, the Irish investigative-intelligence platform — the same category, selling to law enforcement, intelligence and financial-crime teams. **Solutions engineer: managed customer accounts and built the customisations each client needed to make the platform fit their process.** Spoke with hundreds of analysts across conference talks and deployments. Now a sales solutions engineer at UdeMy across large-enterprise and strategic accounts — the same discipline at enterprise scale. 10+ years in data engineering and AI systems.

Full-time on close.

Conor Bowles

CTO

Architect of the case model. The approach the platform rests on is his: extract entities and events from evidence with an LLM at ingestion, resolve them into a graph, and query that concrete structure afterwards — so an answer is a traversal over modelled evidence, not a retrieval over text. He wrote the first commit — LLM client, graph client and ingestion pipeline — and 240+ since: case-level isolation enforced in the database, entity resolution, the split between verified facts and AI inference, and the performance work that keeps very large cases traversable.

Full-time on close.

Alexandra Solórzano

CPO

10+ years as a licensed private investigator on federal criminal defence and financial fraud casework — and the product's direction, daily. **Multiple calls a week; the backlog is generated from live matters as she works them,** which is why development has never run on a roadmap invented in advance. A grounded visionary: every feature traces to a real case, and she has run federal cases to completion inside the platform. US-based.

The domain authority — and the customer the product was built for.

Seven years building investigative intelligence software for enterprise buyers, and a decade running federal casework as the kind of investigator who has to live with the result. **Our CPO has run federal cases to completion inside this platform.**

STATUS AND TRACTION

Loupe is in production use on real federal casework.

Built inside **Owl Consultancy**, a working US private-investigations practice, and used there to carry federal cases **to completion** — multi-gigabyte phone extractions, tens of thousands of curated financial transactions, full document corpora. **Three attorneys on those cases have been given access to it.** Attorneys have retained the practice specifically because of the platform — unprompted, before any marketing existed.

This is not a prototype seeking a first user. It is an application that works, in the hands of the investigators it was built for, on cases with real stakes.

First external pilot in motion.

A US private-investigations firm led by a former FBI undercover agent — clients across Miami, New York, California and the Midwest — has agreed to send its next large-document matter through Loupe as a paid trial at cost, and to put a mutual MSA in front of us. Sourced through casework, not marketing.

Built so far, verifiable

- ~180,000 lines across a React console, a FastAPI case API and a separate asynchronous evidence engine
- 913 commits since December 2025 — two primary engineers
- Neo4j (case graph) · Postgres (cases, auth, audit) · ChromaDB (retrieval) · Redis (jobs)
- Three AI providers behind one routing policy — the model is a swappable component
- Single-tenant deployment: one isolated stack per customer

Release discipline

Onboarding an external customer's evidence runs behind a **12-stage security-gated release plan. No external case data enters the platform until the security, isolation, backup and legal stages pass.**

Roadmap: a 13-epic board takes Loupe from a platform that works for the practice it was built in to one that scales across many. Depth, not existence.

THE ASK

€100k. What it buys: the first paying customer outside the practice that built it.

The platform works and has closed real federal cases. What stands between it and external revenue is the **security gate** — the isolation, backup and legal stages that let another firm's privileged evidence enter the system — and the pilot itself.

Use of funds

- **Two founders full-time** — clearing the security gate and running the agreed pilot to a signed contract. The pilot candidate is already secured.
- **The security evidence a firm's counsel asks for** — Cyber Essentials, an independent penetration test and a written security package, plus starting the SOC 2 observation window. Cheap and fast; full certification is bought when a deal requires it.
- **Single-tenant infrastructure** for the first external customers — one instance each, so it scales with customers rather than ahead of them.
- **Processing and inference** across pilot matters. The per-matter cost model is what the pilot establishes.

What we want from NDRC beyond capital

- **Go-to-market guidance.** The channel is identified — bar associations and professional bodies — and pricing is per-matter rather than per-seat. Both need pressure-testing by people who have built a repeatable motion.
- Dublin base at Dogpatch Labs
- EIR time on pricing structure and enterprise sales motion
- Introductions into Irish and UK professional-services buyers — ACFE Ireland, Chartered Accountants Ireland forensic group, Law Society criminal law committee

We are not asking anyone to fund the invention of a product. We are asking them to fund the distance between works and sold.