



Find the signal in everything.

An investigation platform for fraud and criminal casework.

Dublin, Ireland.



THE PROBLEM

Modern casework is a data problem wearing a legal costume.

A single serious-fraud or criminal case now arrives as:

- thousands of documents
- multi-gigabyte phone extractions, hundreds of thousands of events
- bank records spanning tens of thousands of transactions
- hours of recorded audio

The tools investigators have are **viewers**. One phone at a time. One PDF at a time. One spreadsheet at a time.

The connections — the same person in a witness statement, a WhatsApp thread and a beneficiary field; the meeting two streets from the cash withdrawal — live in investigators' heads and on whiteboards.

They don't scale, they don't survive staff turnover, and they can't be handed to a court.

We didn't start with a product idea. We started with a case we couldn't work.

We tried to buy this. Building it was the second choice, not the first.

Three curves crossed.

01

Digital evidence volume exploded. A phone extraction now appears in nearly every case.

02

LLMs became good enough to read evidence at scale.

03

Bare LLMs remain unusable in evidence work — no provenance, no audit, no confidentiality.

The gap between (2) and (3) is exactly where Loupe sits:

the trust layer that makes AI admissible in investigative work.

Only ~25% of anti-fraud programmes use AI/ML today — and only **~6% of fraud examiners are confident explaining how their AI reaches its conclusions.**

ACFE 2026 anti-fraud technology benchmarking

PRODUCT

Evidence in. One connected model out.

Phone extractions and call logs, documents, images, audio, financial records — everything a case arrives with — resolved at ingestion into **one model**, where a call, a transfer and a line in a subpoena return are the same kind of citable object.

ONE MODEL, MANY PERSPECTIVES

Financial explorer

Multi-phone view

Timeline

Graph

Agent query

SCREENSHOT — A LOUPE

Passages, entities and events from several sources,
bound into one evidenced finding.

SCREENSHOT — CITED AI ANSWER, MID-CLICK

Opening the source document
at the cited page.

A Loupe is a bonded collection — the passages, entities, events and transactions that establish one thing, each carrying the quote, page and file it came from.

Investigators work the perspectives, find what matters, and bind it into a Loupe. **Loupes move into the case, and the case is built out of them** — conclusion, justification and evidence, assembled as the investigation runs rather than written up at the end.

DIFFERENTIATION

Why this isn't ChatGPT with files.

	A straight LLM	Loupe
Where the case lives	A context window, forgotten after	A persistent case knowledge graph — outlives every conversation
Scale	Hundreds of pages at best	Entire corpora: thousands of documents, multi-gigabyte phone reports
What the AI does	Answers in prose, one conversation at a time	Cites every claim to the page — and runs tools: builds Loupes, groupings and outputs from plain instruction
Confidentiality	Multi-tenant cloud	Single-tenant: evidence never leaves the customer's instance
Human judgment	Lives outside the tool — lost between sessions	Highlights, significance and Loupes — bonded collections of evidence — are durable objects in the case model
Model dependence	The model is the product	Models are interchangeable components. The moat is the evidence layer.

A straight LLM gives you a well-written opinion about the fraction of the evidence that fits in its context window. Loupe gives you a complete, permanent, cross-referenced model of all of it — where **every claim carries its receipt.**

MARKET AND LEAD CUSTOMER

Every feature was requested by a working investigator on a live case.

Our CPO is our design partner. Development has been customer-led from the first commit — not a roadmap we invented and then validated, but a backlog generated by real federal casework. Cases have been worked to completion in the platform, and **three attorneys on those cases have been given access to it.**

Why this tier is empty — and stays empty.

Enterprise investigative-intelligence platforms are sold with configuration and support attached — the right model for agencies and large institutions, who arrive with budget, procurement and implementation resource.

The tier below has the same casework complexity and the same evidence volumes — and none of that budget. **Loupe is built for it directly: usable out of the box, by a practice with no implementation resource at all.**

We did not set out to build a platform. We went looking for one, tried several, and none could work the case.

Beachhead

Boutique forensic and criminal-defence practices — 10–50 people, Big-Four-class casework without Big-Four tooling budgets.

Expansion

Financial-crime & fraud units · law enforcement & prosecutors · law firms & disputes teams · corporate investigations & compliance.

Channel

This buyer is organised and reachable. ACFE has ~95k members and ~60k CFEs worldwide, with chapters, conferences and directories — including an Ireland chapter since 2016. The defence bar is the same shape: a competitor reached national presence through a national association and one state bar.

Why we know the gap is real

We tried a number of platforms on live federal casework before we built anything. None of them could work the case.

GROWTH PATH

Ten million in ARR is 104 firms and under 2% of one market.

Priced at a blended \$12,000 per matter. A boutique running eight evidence-heavy matters a year is worth \$96,000 annually; an active practice at twenty is worth \$240,000. These figures use the conservative rate throughout.

STAGE 1 • PROOF

Owl & Ireland

\$0.5M ARR

5 customers

42 matters / yr

References and a validated price. Not a revenue stage.

STAGE 2 • THE ENGINE

United States

\$5M ARR

52 customers

417 matters / yr

0.3–0.9% of 48,000–125,000 evidence-heavy matters a year.

STAGE 3 • SECOND GEOGRAPHY

+ United Kingdom

\$10M ARR

104 customers

833 matters / yr

80,200 open Crown Court cases, rising. Same buyer, same channel.

STAGE 4 • EXPANSION

+ Europe

\$25M ARR

260 customers

2,083 matters / yr

\$2.41B forensics, \$6.81B legal tech. Partner-led per jurisdiction.

TEAM

We didn't research this market. We worked in it — from both sides.

Neil Byrne

CEO

Seven years at Siren, the Irish investigative-intelligence platform — the same category, selling to law enforcement, intelligence and financial-crime teams. **Solutions engineer: managed customer accounts and built the customisations each client needed to make the platform fit their process.** Spoke with hundreds of analysts across conference talks and deployments. Now a sales solutions engineer at Udemy across large-enterprise and strategic accounts — the same discipline at enterprise scale. 10+ years in data engineering and AI systems.

Full-time on close.

Conor Bowles

CTO

Architect of the case model. The approach the platform rests on is his: extract entities and events from evidence with an LLM at ingestion, resolve them into a graph, and query that concrete structure afterwards — so an answer is a traversal over modelled evidence, not a retrieval over text. He wrote the first commit and has led engineering since. **Drives the enterprise platform**, while features are proven on live casework first and promoted once they mature — so nothing reaches the deployable product until it has survived a real matter.

Full-time on close.

Alexandra Solórzano

CPO

10+ years as a licensed private investigator on federal criminal defence and financial fraud casework — and the product's direction, daily. **Multiple calls a week; the backlog is generated from live matters as she works them**, which is why development has never run on a roadmap invented in advance. A grounded visionary: every feature traces to a real case, and she has run federal cases to completion inside the platform. US-based.

The domain authority — and the customer the product was built for.

Seven years building investigative intelligence software for enterprise buyers, and a decade running federal casework as the kind of investigator who has to live with the result. **Our CPO has run federal cases to completion inside this platform.**

STATUS AND TRACTION

Loupe is in production use on real federal casework.

Built inside **Owl Consultancy**, a working US private-investigations practice, and used there to carry federal cases **to completion** — multi-gigabyte phone extractions, tens of thousands of curated financial transactions, full document corpora. **Three attorneys on those cases have been given access to it.** Attorneys have retained the practice specifically because of the platform — unprompted, before any marketing existed.

This is not a prototype seeking a first user. It is an application that works, in the hands of the investigators it was built for, on cases with real stakes.

First external pilot in motion.

A US private-investigations firm led by a former FBI undercover agent — clients across Miami, New York, California and the Midwest — has agreed to send its next large-document matter through Loupe as a paid trial at cost, and to put a mutual MSA in front of us. Sourced through casework, not marketing.

Built so far, verifiable

- ~180,000 lines across a React console, a FastAPI case API and a separate asynchronous evidence engine
- 913 commits since December 2025 — two primary engineers
- Neo4j (case graph) · Postgres (cases, auth, audit) · ChromaDB (retrieval) · Redis (jobs)
- Three AI providers behind one routing policy — the model is a swappable component
- Single-tenant deployment: one isolated stack per customer

Release discipline

Onboarding an external customer's evidence runs behind a **12-stage security-gated release plan. No external case data enters the platform until the security, isolation, backup and legal stages pass.**

Roadmap: a 13-epic board takes Loupe from a platform that works for the practice it was built in to one that scales across many. Depth, not existence.

THE ASK

€100k. What it buys: the first paying customer outside the practice that built it.

The platform works and has closed real federal cases. What stands between it and external revenue is the **security gate** — the isolation, backup and legal stages that let another firm's privileged evidence enter the system — and the pilot itself.

Use of funds

- **Two founders full-time** — clearing the security gate and running the agreed pilot to a signed contract. The pilot candidate is already secured.
- **The security evidence a firm's counsel asks for** — Cyber Essentials, an independent penetration test and a written security package, plus starting the SOC 2 observation window. Cheap and fast; full certification is bought when a deal requires it.
- **Single-tenant infrastructure** for the first external customers — one instance each, so it scales with customers rather than ahead of them.
- **Processing and inference** across pilot matters. The per-matter cost model is what the pilot establishes.

What we want from NDRC beyond capital

- **Go-to-market guidance.** The channel is identified — bar associations and professional bodies — and pricing is per-matter rather than per-seat. Both need pressure-testing by people who have built a repeatable motion.
- Dublin base at Dogpatch Labs
- EIR time on pricing structure and enterprise sales motion
- Introductions into Irish and UK professional-services buyers — ACFE Ireland, Chartered Accountants Ireland forensic group, Law Society criminal law committee

We are not asking anyone to fund the invention of a product. We are asking them to fund the distance between works and sold.